



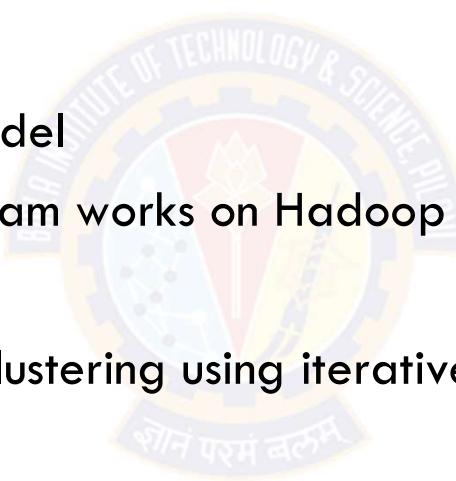
BITS Pilani
Pilani | Dubai | Goa | Hyderabad

DSECL ZG 522: Big Data Systems Session 7 - Distributed Programming

Janardhanan PS
janardhanan.ps@wilp.bits-pilani.ac.in

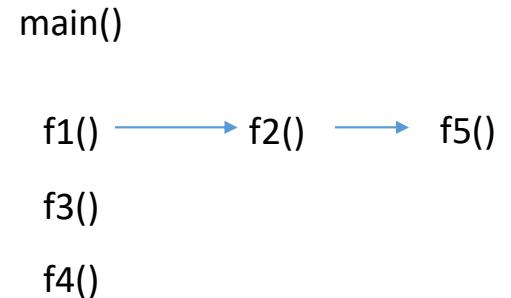
Topics for today

- **Top down design**
- Types of parallelism
- MapReduce programming model
- See how a map reduce program works on Hadoop
- Iterative MapReduce
- Hands on demo of K-Means clustering using iterative MapReduce



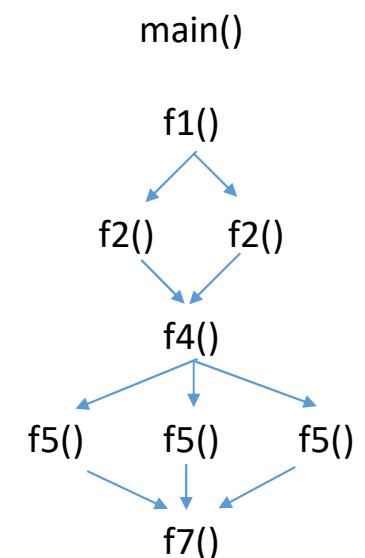
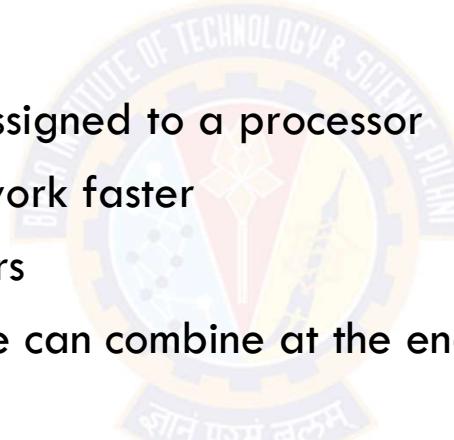
Top down design - Sequential context

- In the context of a sequential program
 - Divide and conquer
 - It is easier to divide a problem into sub-problems and execute them one by one on single CPU
 - A sub-problem definition may be left to the programmer in a sequential programming context



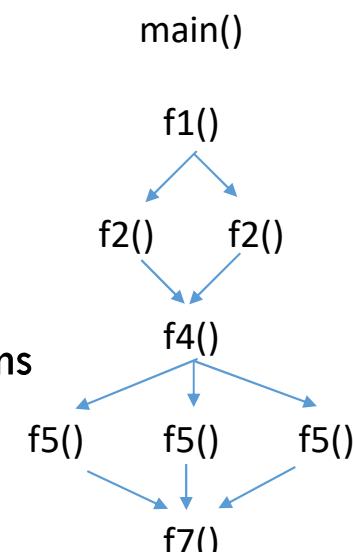
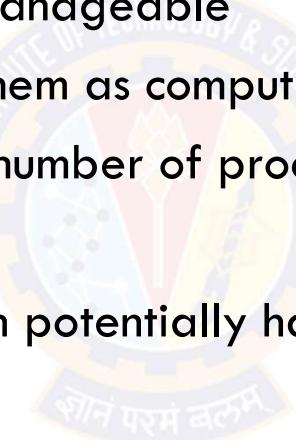
Top down design – Parallel context

- In the context of a parallel program, we cannot decompose the problem into sub-problems in anyway the programmer chooses to
- Need to think about
 - Each sub-problem needs to be assigned to a processor
 - Goal is to get the program work faster
 - Utilise all available processors
 - Divide the problem only when we can combine at the end into the final answer
 - Need to decide where to do the combination
 - Is there any parallelism in combination or is it sequential or trivial



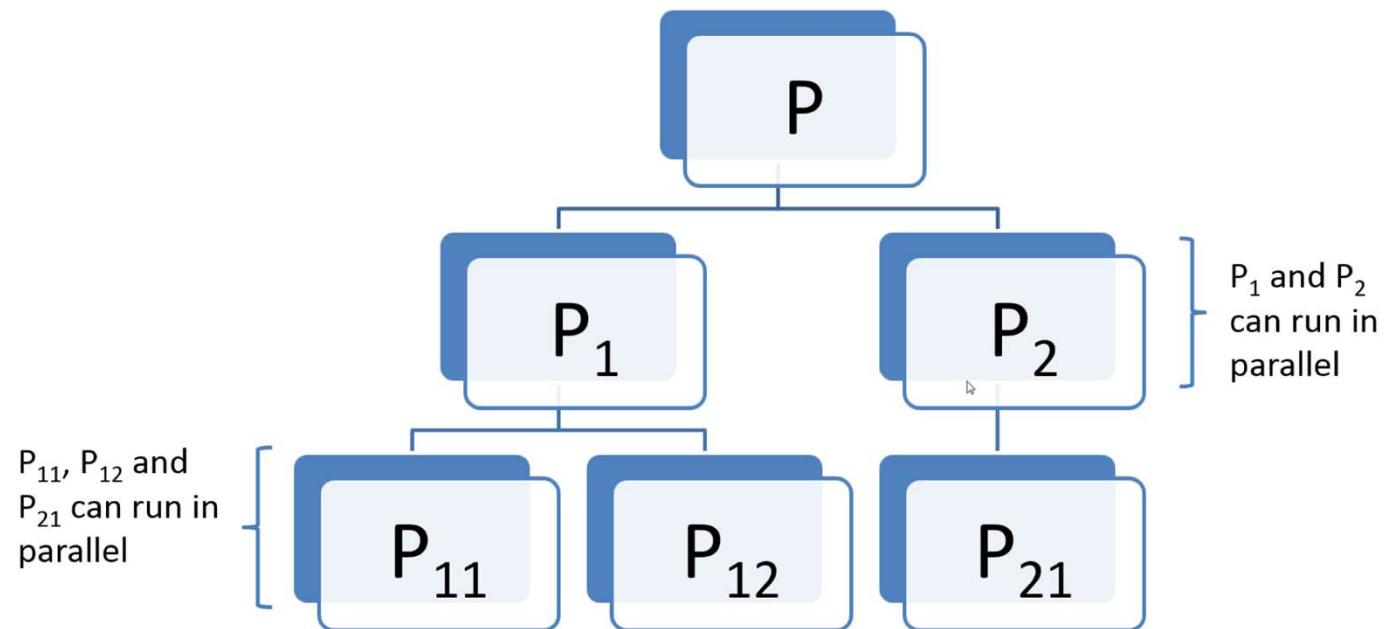
Deciding on number of sub-problems

- In conventional top-down design for sequential systems
 - Keep number of sub-problems manageable
 - Because need to keep track of them as computation progresses
- In parallel systems, it is dictated by number of processors
 - Processor utilization is the key
 - If there are N processors, we can potentially have N sub-problems



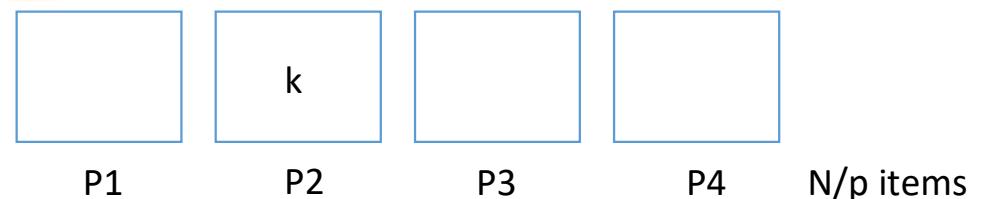
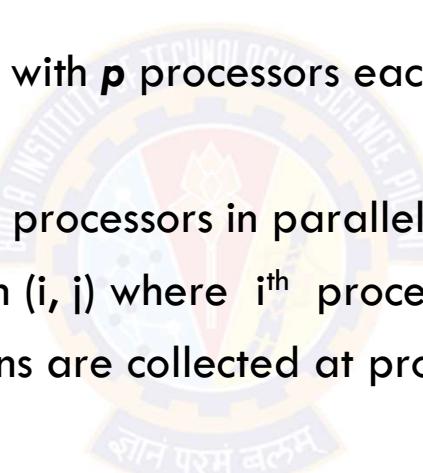
Top-down design

- At each level, problems need to run in parallel



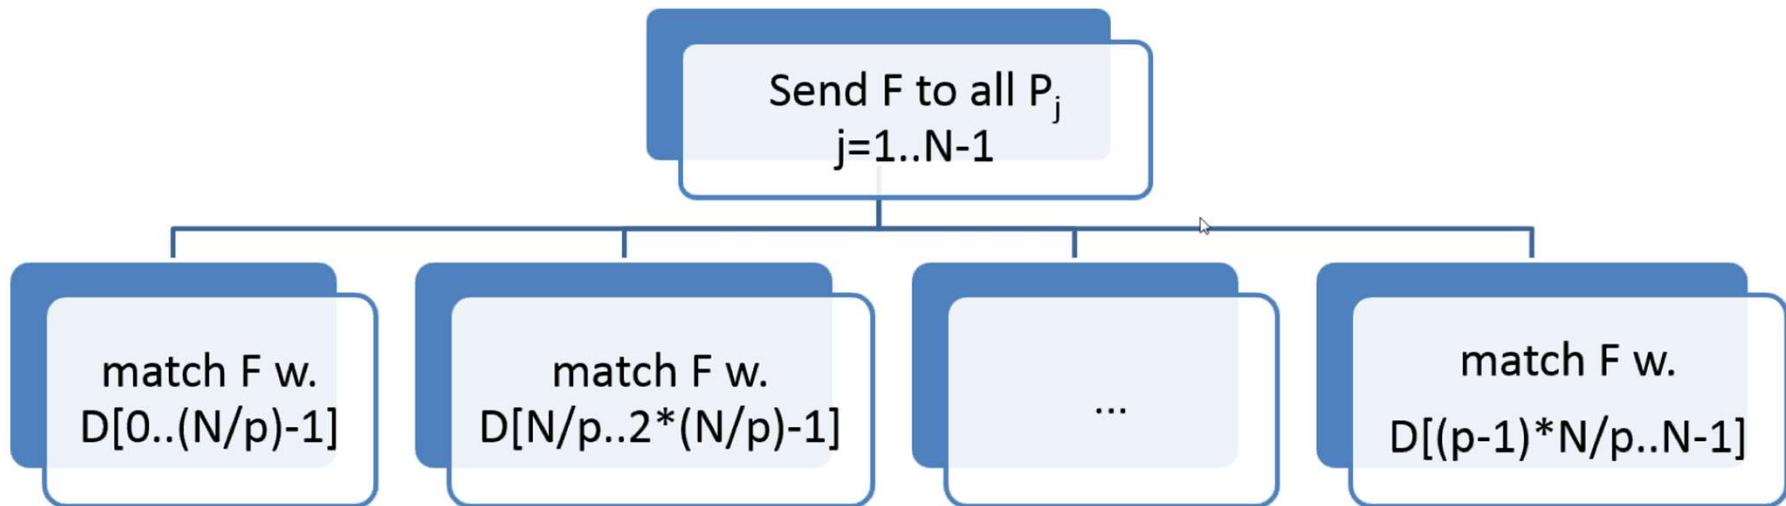
Example 1 - Keyword search in list

- Problem:
 - Search for a key k in a sorted list L_s of size N
- Data:
 - L_s is stored in a distributed system with p processors each storing N/p items
- Solution:
 - Run binary search in each of the p processors in parallel
 - Whichever processor finds k return (i, j) where i^{th} processor has found key in j^{th} position
 - Combination: One or more positions are collected at processor 0
- Speedup: p
- Time complexity: $O(\log(N/p))$



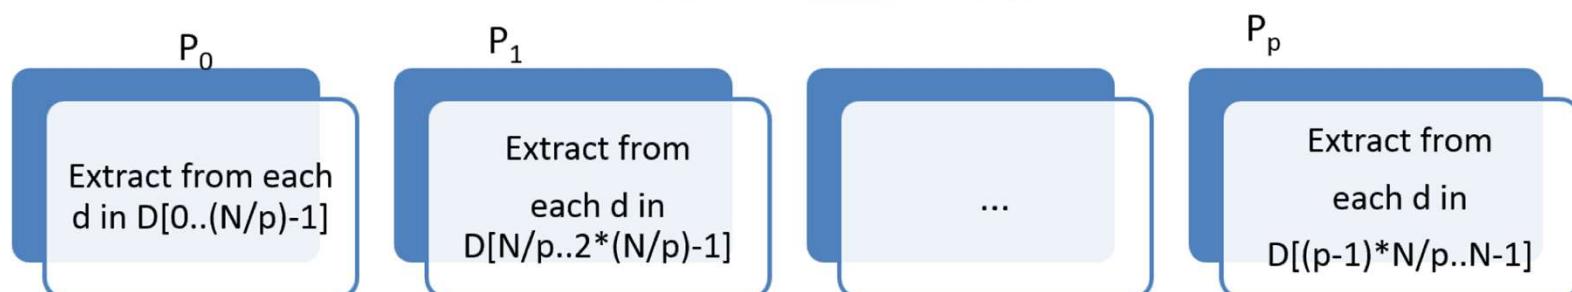
Example 2 - Fingerprint matching

- Find matches for a fingerprint F in a database of D prints
- Set of D prints is partitioned and evenly stored in a distributed database
- Partitioning is an infrequent activity - only when many new entries in database
- Search is the frequent activity
- Speed up p
- Time complexity $O(N/p)$ given sequential search in every partition



Example 3: Document search

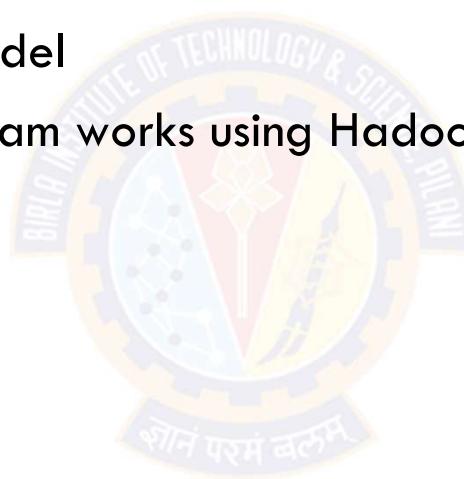
- Find keywords from each document d in a distributed document collection D



P_0 : Collect keywords from all processors

Topics for today

- Top-down design
- **Types of parallelism**
- MapReduce programming model
- See how a map reduce program works using Hadoop
- Iterative MapReduce



Types of Parallelism

1. Data Parallelism
2. Tree Parallelism
3. Task Parallelism
4. Request Parallelism



1. Data parallel execution model

- Data is partitioned to multiple nodes / processors
 - Try to make partitions equal or balanced
- All processors execute the same code in parallel
 - For homogenous nodes and equal amount of work, the utilization will be close to 100%
 - Execution time overhead is minimal
 - Unbalanced data size / work or heterogenous nodes will lead to higher execution time



Where data parallelism is not possible

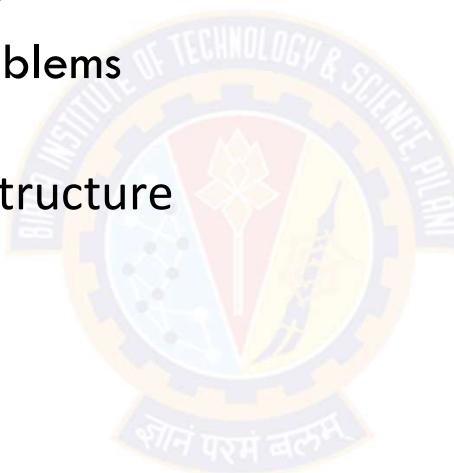
- There are problems where you cannot divide the work
 1. equally
 2. independently to proceed in parallel
- QuickSort(L_s, N)
 - All N items in L_s have to be in memory of processor 0
 - Time Complexity
 - Best case - $O(n \log(n))$
 - Worst case - $O(n^2)$



[Sorting \(Bubble, Selection, Insertion, Merge, Quick, Counting, Radix\) - VisuAlgo](#)

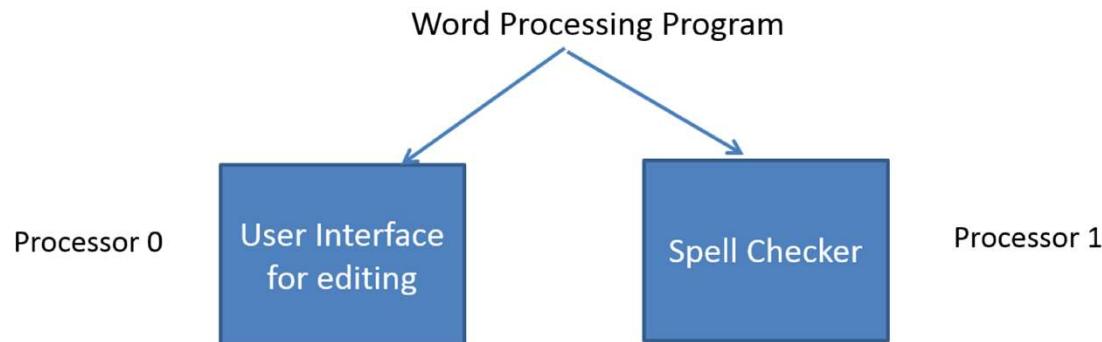
2. Tree parallelism

- Dynamic version of divide and conquer - partitions are done dynamically
- Division of problem into sub-problems happens at execution time
 - Sub-problem is identical in structure to the larger problem
 - What is the division size ?



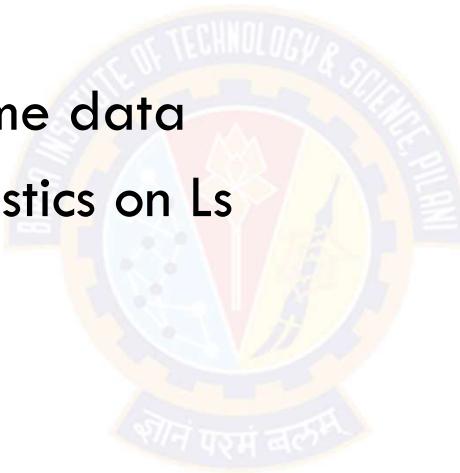
3. Task parallelism - Example 1 : Word processor

- Parallel tasks that work on the same data
 - Unlike data and tree parallelism, data doesn't need to be divided, the Task gets divided into sub-tasks
 - May work on same data instance, else need to make data copies and keep them in sync
- If on multiple core, different threads can execute tasks in parallel accessing same data instance in memory



Task parallelism - Example 2 : Independent statistics

- Given a list Ls of numeric values find its mean, median and mode
- Solution
 - Independent tasks on same data
 - Each task can find a statistics on Ls
 - Run tasks in parallel



Task parallelism summary

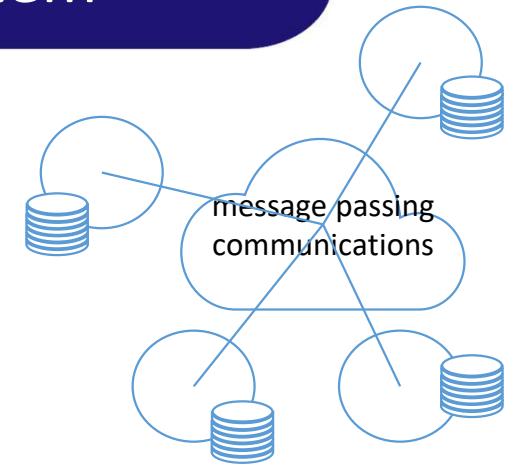
- Identify sub-tasks based on functionality with no common function
 - In Tree and Data parallel, the tasks are identical functions
- Sub-tasks are not identified based on data
- Independent sub-tasks are executed in parallel
- Sub-tasks are often limited and known statically in advance
 - We know in a word processor, what are the sub-tasks
 - We know in statistical analysis, what functions we will run in advance
 - So limited parallelism scope - not scalable with more resources
 - In data or tree parallelism, we can potentially get more parallelism with more data
 - more scalable with more resources at same time interval **BigData**

4. Request parallelism

- Problem
 - Scalable execution of independent tasks in parallel
 - Execute same code but in many parallel instances
- Solution
 - On arrival, each request is serviced in parallel along with other existing tasks servicing prior requests
 - Could be processing same or fixed data
 - Request-reply pairs are independent of each other serviced by a different thread or process in the backend
 - There could be some application specific backend dependency, e.g. GET and POST on same data item
- Systems Fit
 - Servers in client-server models
 - e.g. email server, HTTP web-server, cloud services with API interface, file / storage servers
 - Microservices
 - Socket programming
- Scalability metrics : Requests / time (throughput)

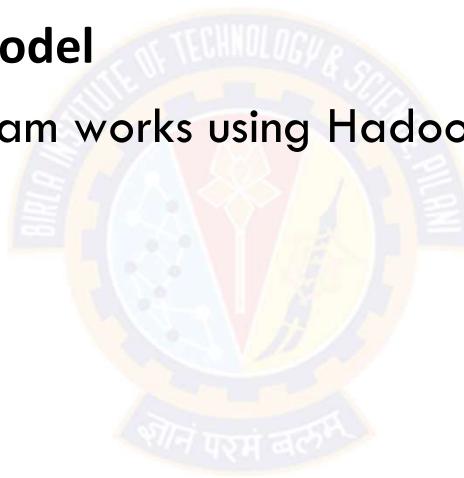
What happens in a loosely coupled distributed system

- Divide
 - No shared memory
 - Memory / Storage is on separate nodes
 - So, any exchange of data or coordination between tasks is via message passing
 - Divide the problem in a way that computation task can run on local data
- Conquer / Merge
 - In shared memory merge it is simpler with each process writing into a memory location
 - In distributed
 - Need to collect data from the different nodes
 - In search example, it is a simpler merge to just collect result - so low cost
 - In quick sort, it is simple append whether writing in place for shared memory or sending a message
 - Sometimes merges may become sequential
 - e.g. k-means - in each iteration (a) guess clusters in parallel to improve the clusters but (2) checking if we have found right clusters is sequential



Topics for today

- Top down design
- Types of parallelism
- **MapReduce programming model**
- See how a map reduce program works using Hadoop
- Iterative MapReduce

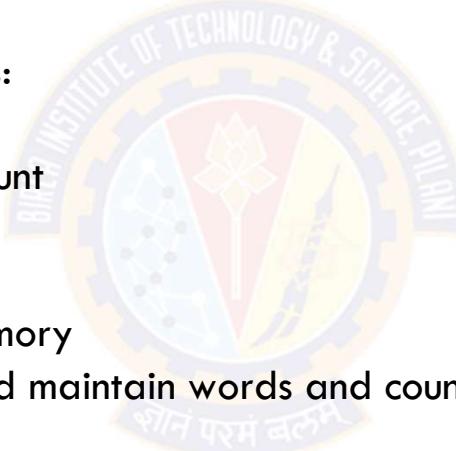


MapReduce origins

- Created in Google to run on GFS
- Open-source version created as Apache Hadoop
- Perform maps/reduces on data using many machines
 - The framework takes care of distributing the data and managing fault tolerance
 - You just write code to map one element and reduce elements to a combined result
 - Separates how to do recursive divide-and-conquer from what computation to perform
 - Old idea in higher-order functional programming transferred to large-scale distributed computing
 - Complementary approach to database declarative queries
 - In SQL, you don't actually write the low-level query execution code
 - Programmer needs to focus just on **map** and **reduce** logic and rest of the work is done by the map-reduce framework.
 - So restricted programming interface to the system to let the system do the distribution of work, job tracking, fault tolerance etc.

MapReduce Evolution

- We have a large text file of words, one word in a line
- We need to count the number of times each distinct word appears in the file
- Sample application: analyze web server logs to find popular URLs
- Word Count Solution Design Options:
 - 1: Entire file fits in memory
 - Read file into memory and count
 - 2: File too large for memory, but all $\langle \text{word}, \text{count} \rangle$ pairs fit in memory
 - Read file line by line and maintain words and counts in memory
 - 3: File on disk, too many distinct words to fit in memory
 - sort words.txt | uniq -c**



Solution for multiple large files on disk

To make it slightly harder, suppose we have a large corpus of documents

Count the number of times each distinct word occurs in the corpus

- `words (docs/*) | sort | uniq -c`

where `words` takes a file in a folder and outputs the words in it, one to a line

- The above captures the essence of MapReduce
- Great thing is it is naturally parallelizable

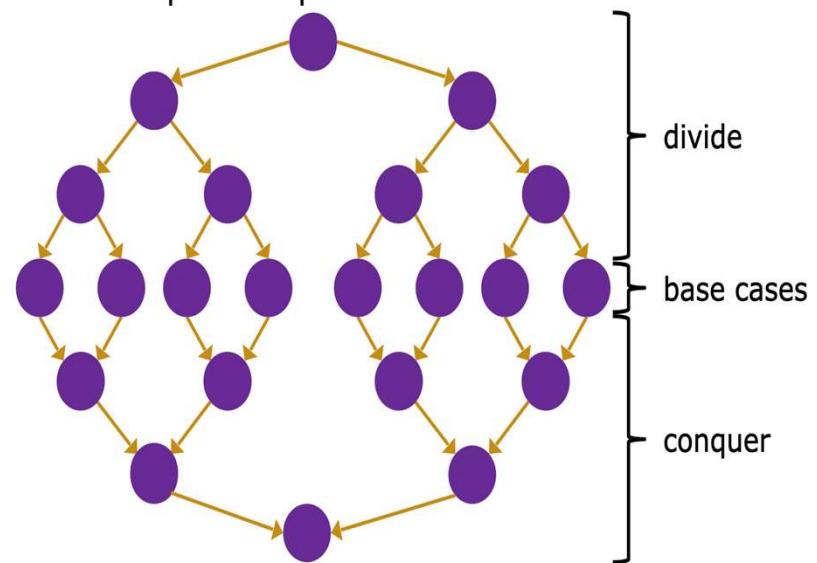
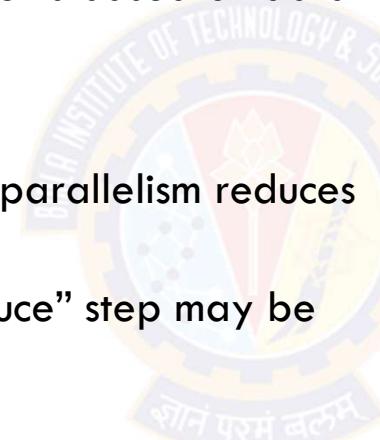
MapReduce is conceptually like a UNIX pipeline

- One function (Map) processes data
 - Output is input to another function (Reduce)
- ```
cat words.txt | sort | uniq -c | cat > file
input | map | shuffle | reduce | output
```

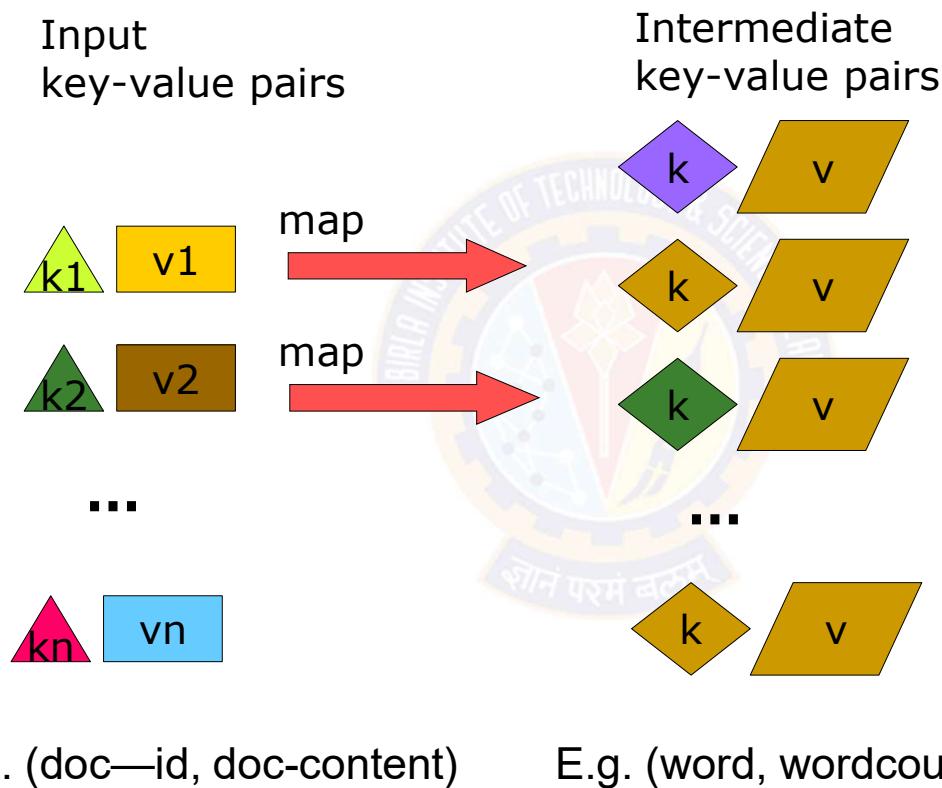
- Developer specifies two functions:
  - `map()` - User code
  - `reduce()` - User code
- Rest of the job is done by the MapReduce framework
- Tune the configuration parameters of the MapReduce framework for performance

# MapReduce in terms of Data and Tree parallelism

- Map
  - Data parallelism
  - Divide a problem into sub-problems based on data
- Reduce
  - Inverse tree parallelism
  - With every merge / reduce the parallelism reduces until we get one result
  - Depending on the problem “reduce” step may be simple or sequential



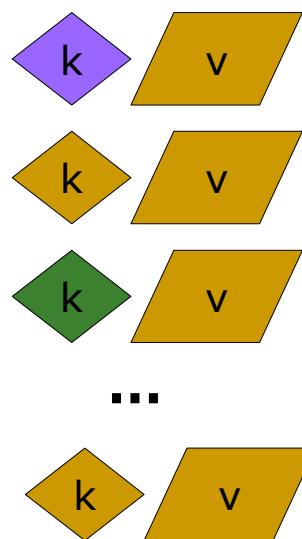
# MapReduce: The Map Step



Adapted from Jeff Ullman's course slides<sup>25</sup>

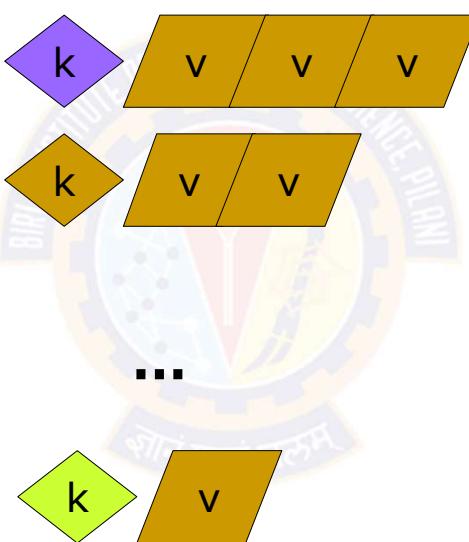
# MapReduce: The Reduce Step

Intermediate key-value pairs



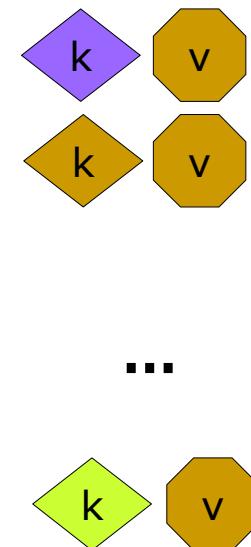
E.g.  
(word, wordcount-in-a-doc)

Key-value groups



(word, list-of-wordcount)  
~ SQL Group by

Output key-value pairs



(word, final-count)  
~ SQL aggregation

Adapted from Jeff Ullman's course slides

## Example: Word Count using MapReduce (Pseudo code)

```
map(key, value):
// key: document name; value: text of document
```

```
 for each word w in value:
```

```
 emit(w, 1)
```

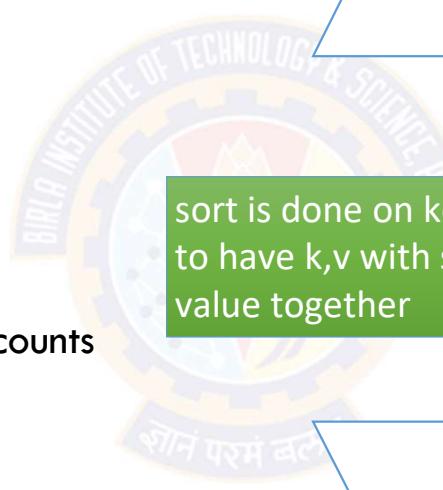
```
reduce(key, values):
// key: a word; value: an iterator over counts
```

```
 result = 0
```

```
 for each count v in values:
```

```
 result += v
```

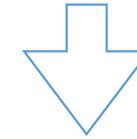
```
 emit(result)
```



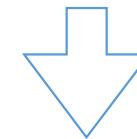
sort is done on keys  
to have k,v with same k  
value together

D1 : the blue ship on blue sea

the, 1    blue, 1    ship, 1    on, 1  
              blue, 1    sea, 1

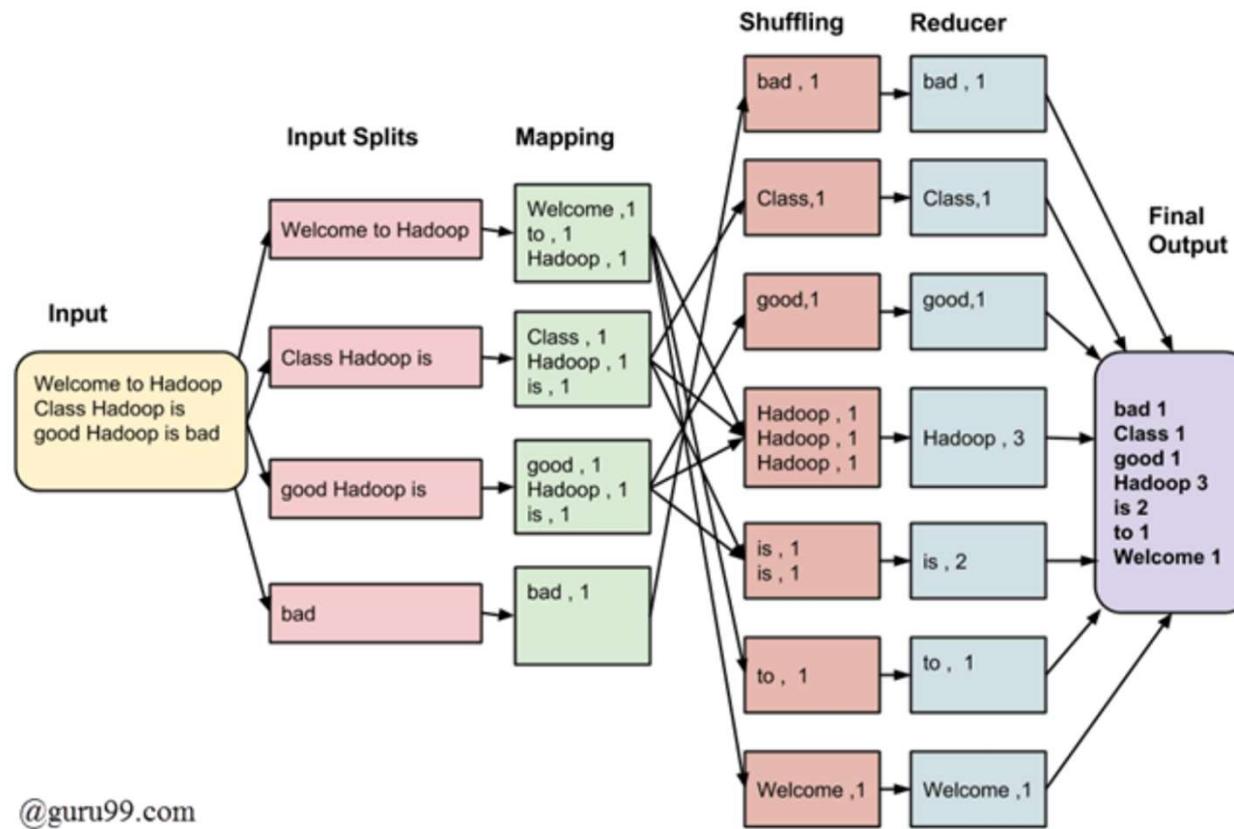


blue, [1,1]    on, 1    sea, 1    ship, 1    the, 1



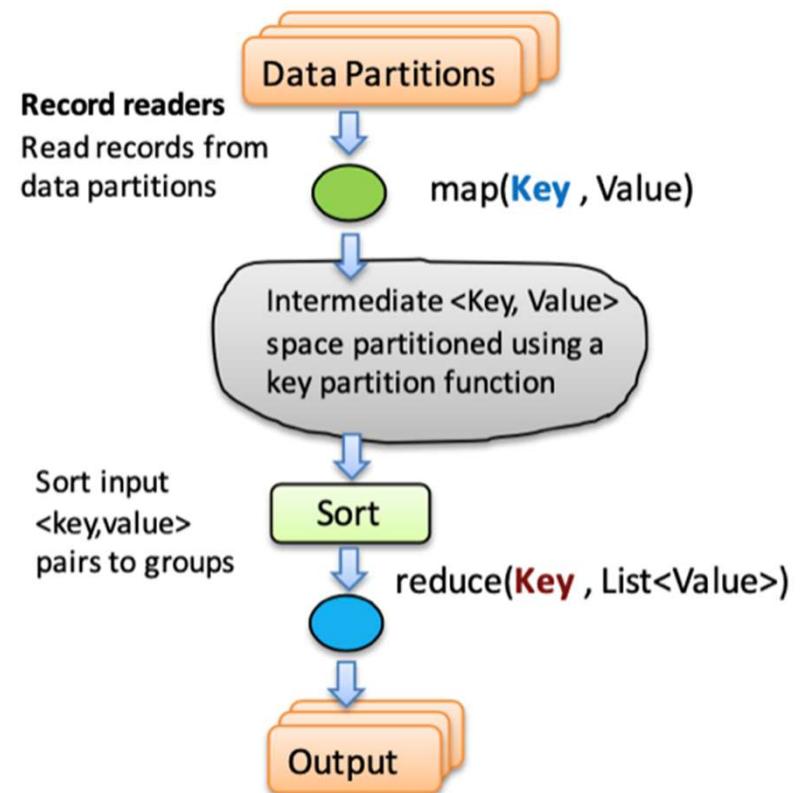
blue, 2    on, 1    sea, 1    ship, 1    the, 1

## Word count (2)



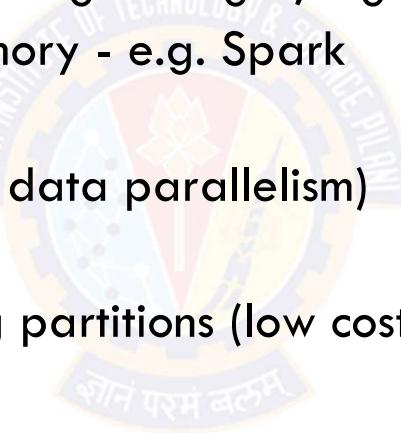
# Formal definition of a MapReduce program

- Input: a set of key/value pairs
- User supplies two functions:
  - $\text{map}(k,v) \rightarrow \text{list}(k_1, v_1)$
  - $\text{reduce}(k_1, \text{list}(v_1)) \rightarrow v_2$
- $(k_1, v_1)$  is an intermediate key/value pair
- Output is the set of  $(k_1, v_2)$  pairs



## When will you use this ?

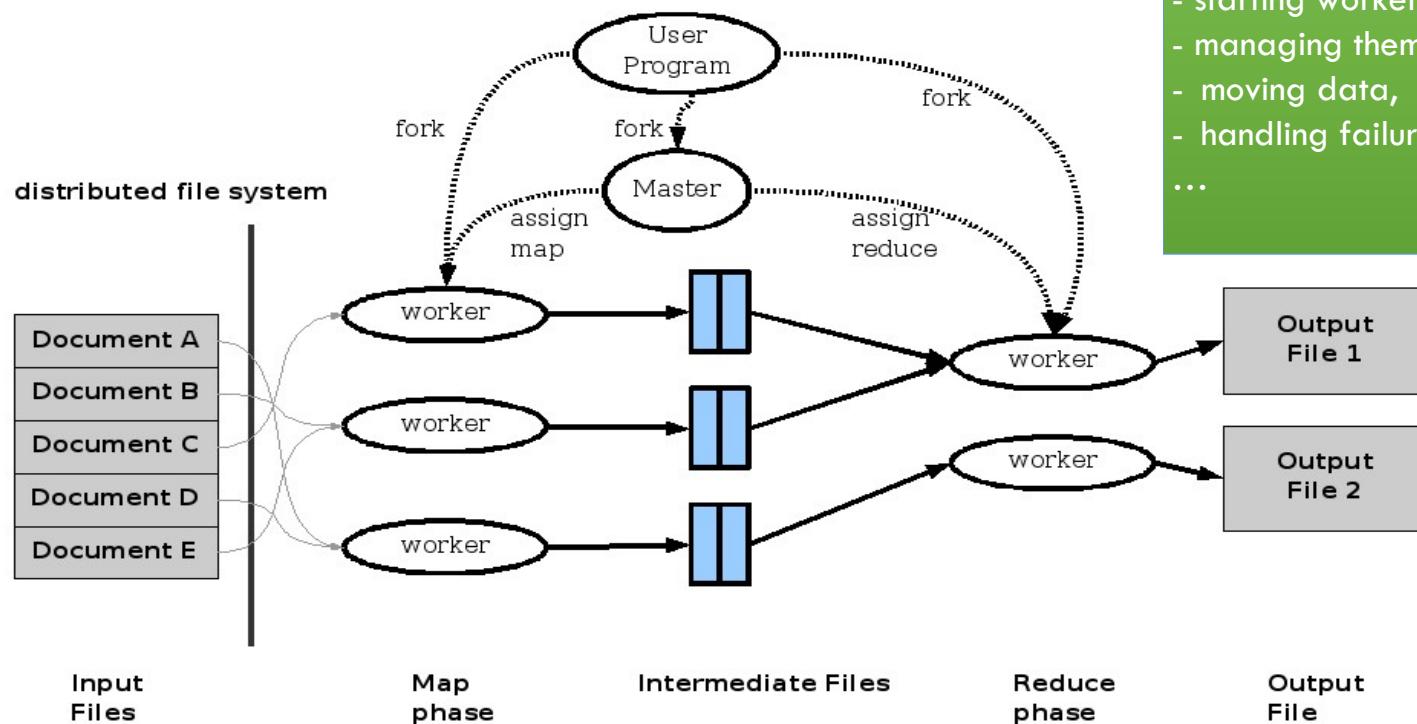
- Huge set of documents that don't fit into memory
  - So, need file-based processing in stages, e.g. Hadoop
  - But can also do this in memory - e.g. Spark
- Lot of data partitioning (high data parallelism)
- Possibly simple merge among partitions (low cost inverse tree parallelism)



# MapReduce: Execution overview

- Data centric design
- Move computation closer to data

- Intermediate results on disk
- Dynamic task scheduling



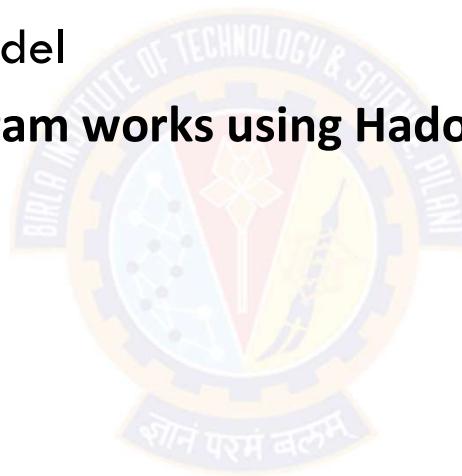
A MapReduce library and runtime does all the work for

- allocating resources,
- starting workers,
- managing them,
- moving data,
- handling failures

...

## Topics for today

- Top down design
- Types of parallelism
- MapReduce programming model
- **See how a map reduce program works using Hadoop**
- Iterative MapReduce



# MapReduce example - sales data processing

Sales by country

| A  | B                | C        | D     | E            | F        | G              | H             | I               | J                | K                | L                |          |          |
|----|------------------|----------|-------|--------------|----------|----------------|---------------|-----------------|------------------|------------------|------------------|----------|----------|
| 1  | Transaction_date | Product  | Price | Payment_Name | City     | State          | Country       | Account_Created | Last_Login       | Latitude         | Longitude        |          |          |
| 2  | 01-02-2009 06:17 | Product1 | 1200  | Mastercard   | carolina | Basildon       | England       | United Kingdom  | 01-02-2009 06:00 | 01-02-2009 06:08 | 51.5             | -1.11667 |          |
| 3  | 01-02-2009 04:53 | Product1 | 1200  | Visa         | Betina   | Parkville      | MO            | United States   | 01-02-2009 04:42 | 01-02-2009 07:49 | 39.195           | -94.6819 |          |
| 4  | 01-02-2009 13:08 | Product1 | 1200  | Mastercard   | Federica | Astoria        | OR            | United States   | 01-01-2009 16:21 | 01-03-2009 12:32 | 46.18806         | -123.83  |          |
| 5  | 01-03-2009 14:44 | Product1 | 1200  | Visa         | Gouya    | Echuca         | Victoria      | Australia       | 9/25/05 21:13    | 01-03-2009 14:22 | -36.1333         | 144.75   |          |
| 6  | 01-04-2009 12:56 | Product2 | 3600  | Visa         | Gerd W.  | Cahaba Heights | AL            | United States   | 11/15/08 15:47   | 01-04-2009 12:45 | 33.52056         | -86.8025 |          |
| 7  | 01-04-2009 13:19 | Product1 | 1200  | Visa         | LAURENCE | Mickleton      | NJ            | United States   | 9/24/08 15:19    | 01-04-2009 13:04 | 39.79            | -75.2381 |          |
| 8  | 01-04-2009 20:11 | Product1 | 1200  | Mastercard   | Fleur    | Peoria         | IL            | United States   | 01-03-2009 09:38 | 01-04-2009 19:45 | 40.69361         | -89.5889 |          |
| 9  | 01-02-2009 20:09 | Product1 | 1200  | Mastercard   | adam     | Martin         | TN            | United States   | 01-02-2009 17:43 | 01-04-2009 20:01 | 36.34333         | -88.8503 |          |
| 10 | 01-04-2009 13:17 | Product1 | 1200  | Mastercard   | Renee    | Eli            | Tel Aviv      | Tel Aviv        | Israel           | 01-04-2009 13:03 | 01-04-2009 22:10 | 32.06667 | 34.76667 |
| 11 | 01-04-2009 14:11 | Product1 | 1200  | Visa         | Aidan    | Chatou         | Ile-de-France | France          | 06-03-2008 04:22 | 01-05-2009 01:17 | 48.88333         | 2.15     |          |
| 12 | 01-05-2009 02:42 | Product1 | 1200  | Diners       | Stacy    | New York       | NY            | United States   | 01-05-2009 02:23 | 01-05-2009 04:59 | 40.71417         | -74.0064 |          |
| 13 | 01-05-2009 05:39 | Product1 | 1200  | Amex         | Heidi    | Eindhoven      | Noord-Brabant | Netherlands     | 01-05-2009 04:55 | 01-05-2009 08:15 | 51.45            | 5.466667 |          |
| 14 | 01-02-2009 09:16 | Product1 | 1200  | Mastercard   | Sean     | Shavano        | FTX           | United States   | 01-02-2009 08:32 | 01-05-2009 09:05 | 29.42389         | -98.4933 |          |
| 15 | 01-05-2009 10:08 | Product1 | 1200  | Visa         | Georgia  | Eagle          | ID            | United States   | 11-11-2008 15:53 | 01-05-2009 10:05 | 43.69556         | -116.353 |          |
| 16 | 01-02-2009 14:18 | Product1 | 1200  | Visa         | Richard  | Riverside      | NJ            | United States   | 12-09-2008 12:07 | 01-05-2009 11:01 | 40.03222         | -74.9578 |          |
| 17 | 01-04-2009 01:05 | Product1 | 1200  | Diners       | Leanne   | Julianstown    | Meath         | Ireland         | 01-04-2009 00:00 | 01-05-2009 13:36 | 53.67722         | -6.31917 |          |
| 18 | 01-05-2009 11:27 | Product1 | 1200  | Visa         | Janet    | Ottawa         | Ontario       | Canada          | 01-05-2009 00:25 | 01-05-2009 10:24 | 45.41667         | 75.7     |          |



count tx by country

<https://www.guru99.com/create-your-first-hadoop-program.html>

# Unravelling a MapReduce program - Driver (1)

```
package SalesCountry; ← package name of jar

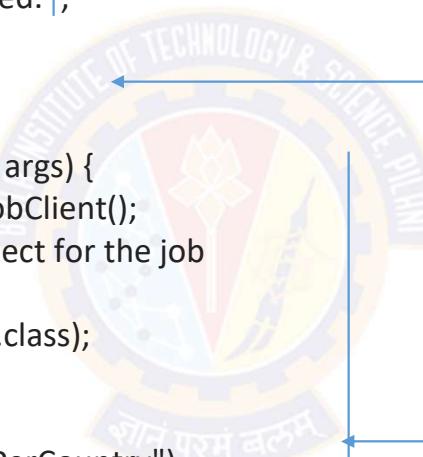
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*; ← hadoop libs

public class SalesCountryDriver { ← driver class that contains main()

 public static void main(String[] args) {
 JobClient my_client = new JobClient();
 // Create a configuration object for the job
 JobConf job_conf = new
 JobConf(SalesCountryDriver.class);

 // Set a name of the Job
 job_conf.setJobName("SalePerCountry");

 // Specify data type of output key and value
 job_conf.setOutputKeyClass(Text.class);
 job_conf.setOutputValueClass(IntWritable.class);
 ...
}
```



Hadoop job with driver class  
with SalesPerCountry as the application  
name

Output data types defined based on  
existing hadoop classes

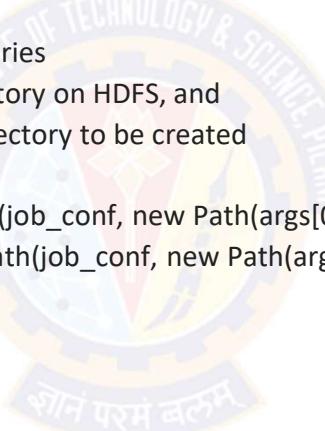
## Unravelling a MapReduce program - Driver (2)

```
...

// Specify names of Mapper and Reducer Class
job_conf.setMapperClass(SalesCountry.SalesMapper.class);
job_conf.setReducerClass(SalesCountry.SalesCountryReducer.class);

// Set input and output directories
// arg[0] = name of input directory on HDFS, and
// arg[1] = name of output directory to be created
// to store the output file.
FileInputFormat.setInputPaths(job_conf, new Path(args[0]));
FileOutputFormat.setOutputPath(job_conf, new Path(args[1]));

my_client.setConf(job_conf);
try {
 // Run the job
 JobClient.runJob(job_conf);
} catch (Exception e) {
 e.printStackTrace();
}
}
```



Mapper and Reducer classes in the jar pkg

Paths to read input and send output

Send job for execution

# Unravelling a MapReduce program - Mapper

```
package SalesCountry;
import java.io.IOException; ← app jar package
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable; ← hadoop libs
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesMapper extends MapReduceBase implements Mapper <LongWritable, Text,
Text, IntWritable> {
 private final static IntWritable one = new IntWritable(1);
 // can add some data structure here for across map() instances
 public void map(LongWritable key, Text value, OutputCollector <Text, IntWritable> output, ←
 Reporter reporter) throws IOException {

 String valueString = value.toString();
 String[] SingleCountryData = valueString.split(",");
 output.collect(new Text(SingleCountryData[7]), one); ←
 } ←
}
```

map function w inputs  
- one or more `<key, value>` pairs  
- output collector  
- ...

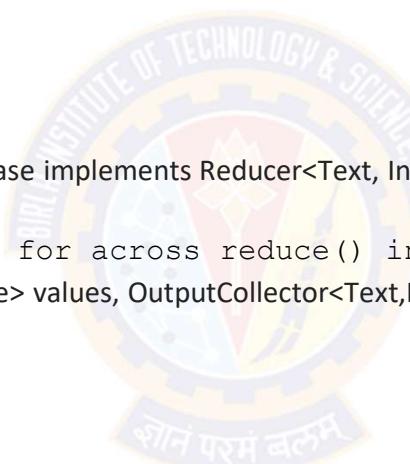
Map logic

`<India, 1> <USA, 1> <India, 1> ....`

# Unravelling a MapReduce program - Reducer

```
package SalesCountry;
import java.io.IOException;
import java.util.*;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesCountryReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text,
IntWritable> {
 // can add some data structure here for across reduce() instances
 public void reduce(Text t_key, Iterator<IntWritable> values, OutputCollector<Text,IntWritable> output,
 Reporter reporter) throws IOException {
 Text key = t_key;
 int frequencyForCountry = 0;
 while (values.hasNext()) {
 // replace type of value with the actual type of our value
 IntWritable value = (IntWritable) values.next();
 frequencyForCountry += value.get();
 }
 output.collect(key, new IntWritable(frequencyForCountry));
 }
}
```



Pkg and includes

reduce function w inputs  
- key and value list, e.g.  
<India, {1,1,1,1}>  
- output collector  
- ...

Reduce logic  
- calculate frequency

For each reduce task, 1 output file created  
Can control #reducer in driver

37

# Running and checking status

```
[[root@centos-s-4vcpu-8gb-blr1-01 source]# ls -l
total 148
-rw-r--r--. 1 root root 44 Apr 18 01:11 Manifest.txt
-rw-r--r--. 1 root root 2966 Apr 18 01:12 ProductSalePerCountry.jar
drwxr-xr-x. 2 root root 96 Apr 19 00:57 SalesCountry
-rw-r--r--. 1 root root 1529 Apr 18 00:57 SalesCountryDriver.java
-rw-r--r--. 1 root root 746 Apr 18 00:56 SalesCountryReducer.java
-rw-r--r--. 1 root root 123637 Jun 6 16:56 SalesJan2009.csv
-rw-r--r--. 1 root root 661 Apr 18 00:56 SalesMapper.java
drwxr-xr-x. 3 root root 157 Jun 6 16:53 units
-rw-r--r--. 1 root root 219 Apr 18 21:10 units.csv
[[root@centos-s-4vcpu-8gb-blr1-01 source]# hadoop jar ProductSalePerCountry.jar /SalesJan2009.csv /output
21/06/06 17:27:37 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
21/06/06 17:27:38 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
21/06/06 17:27:39 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
21/06/06 17:27:39 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute
remedy this.
21/06/06 17:27:40 INFO mapred.FileInputFormat: Total input files to process : 1
21/06/06 17:27:40 INFO mapreduce.JobSubmitter: number of splits:2
21/06/06 17:27:40 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1622978356181_0002
21/06/06 17:27:41 INFO conf.Configuration: resource-types.xml not found
21/06/06 17:27:41 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
21/06/06 17:27:41 INFO resource.ResourceUtils: Adding resource type - name = memory-mb, units = Mi, type = COUNTABLE
21/06/06 17:27:41 INFO resource.ResourceUtils: Adding resource type - name = vcores, units = , type = COUNTABLE
21/06/06 17:27:41 INFO impl.YarnClientImpl: Submitted application application_1622978356181_0002
21/06/06 17:27:41 INFO mapreduce.Job: The url to track the job: http://centos-s-4vcpu-8gb-blr1-01:8088/proxy/application_1622978356181_0002/
21/06/06 17:27:41 INFO mapreduce.Job: Running job: job_1622978356181_0002
21/06/06 17:27:49 INFO mapreduce.Job: Job job_1622978356181_0002 running in uber mode : false
21/06/06 17:27:49 INFO mapreduce.Job: map 0% reduce 0%
21/06/06 17:28:03 INFO mapreduce.Job: map 100% reduce 0%
```

```
[[root@centos-s-4vcpu-8gb-blr1-01 source]# jps
7744 SecondaryNameNode
8357 ResourceManager
8581 NodeManager
13541 Jps
7238 DataNode
10428 MRAppMaster
6910 NameNode
-
```

Logged in as: dr.who



**MapReduce Application application\_1622978356181\_0002**

Active Jobs

| Job ID                 | Name           | State   | Map Progress | Maps Total | Maps Completed | Reduce Progress | Reduces Total | Reduces Completed |
|------------------------|----------------|---------|--------------|------------|----------------|-----------------|---------------|-------------------|
| job_1622978356181_0002 | SalePerCountry | RUNNING | 2            | 2          | 1              | 0               | 0             | 0                 |

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

# MapReduce stats

## File System Counters

```
FILE: Number of bytes read=17747
FILE: Number of bytes written=660936
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=127535
HDFS: Number of bytes written=661
HDFS: Number of read operations=9
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
```

## Job Counters

```
Launched map tasks=2
Launched reduce tasks=1
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)=14658
Total time spent by all reduces in occupied slots (ms)=7011
Total time spent by all map tasks (ms)=14658
Total time spent by all reduce tasks (ms)=7011
Total vcore-milliseconds taken by all map tasks=14658
Total vcore-milliseconds taken by all reduce tasks=7011
Total megabyte-milliseconds taken by all map tasks=15009792
Total megabyte-milliseconds taken by all reduce tasks=7179264
```

## Map-Reduce Framework

```
Map input records=999
Map output records=999
Map output bytes=15743
Map output materialized bytes=17753
Input split bytes=180
Combine input records=0
Combine output records=0
Reduce input groups=58
Reduce shuffle bytes=17753
Reduce input records=999
Reduce output records=58
Spilled Records=1998
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=585
CPU time spent (ms)=2510
Physical memory (bytes) snapshot=772923392
Virtual memory (bytes) snapshot=6393163776
Total committed heap usage (bytes)=527433728
```

## Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

## File Input Format Counters

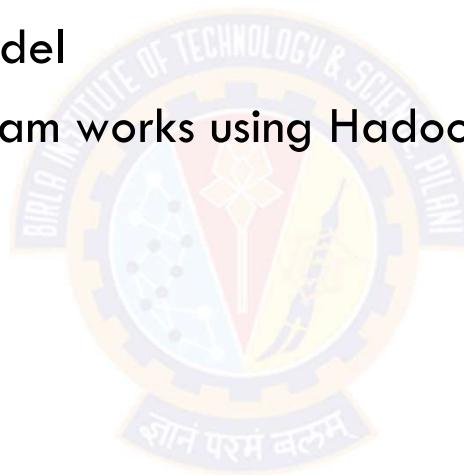
```
Bytes Read=127355
```

## File Output Format Counters

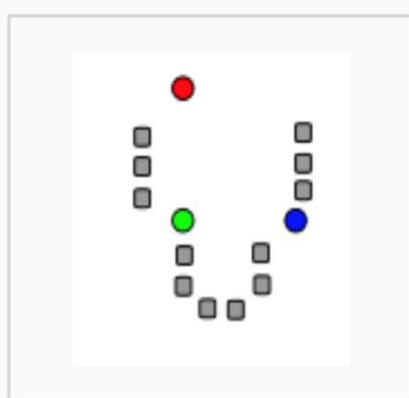
```
Bytes Written=661
```

## Topics for today

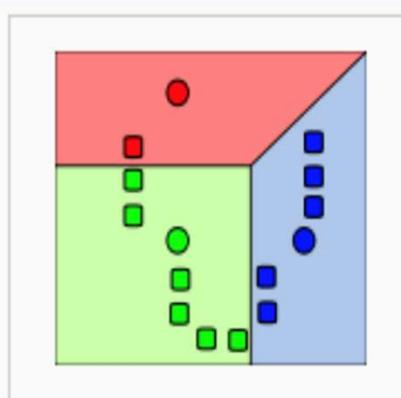
- Top down design
- Types of parallelism
- MapReduce programming model
- See how a map reduce program works using Hadoop
- **K-Means clustering**
- Iterative MapReduce



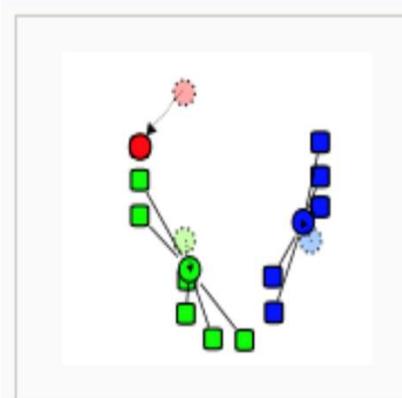
## Example 1: K-means clustering



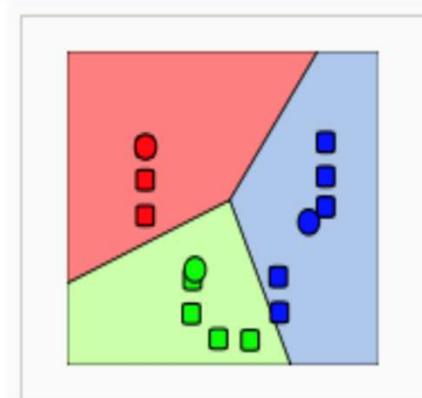
1)  $k$  initial "means" (in this case  $k=3$ ) are randomly selected from the data set (shown in color).



2)  $k$  clusters are created by associating every observation with the nearest mean. The partitions here represent the Voronoi diagram generated by the means.



3) The centroid of each of the  $k$  clusters becomes the new means.

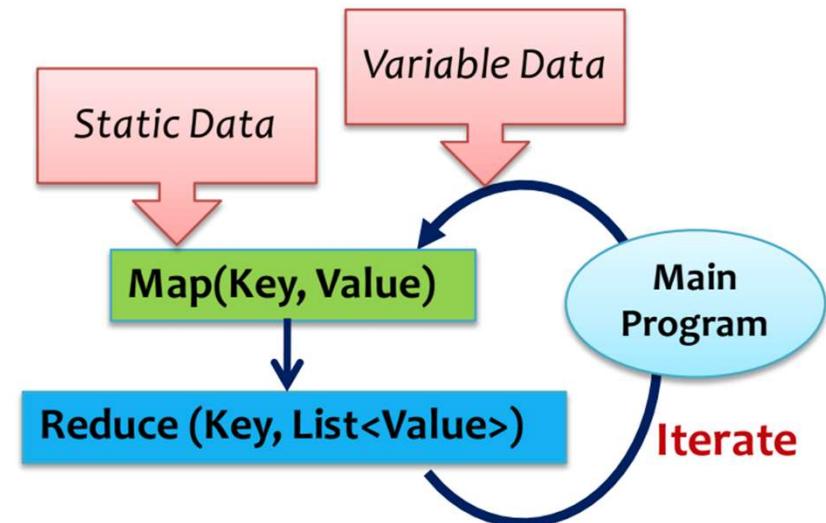
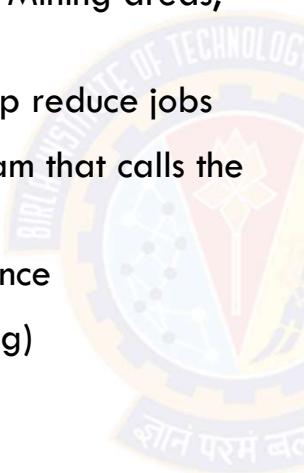


4) Steps 2 and 3 are repeated until convergence has been reached.

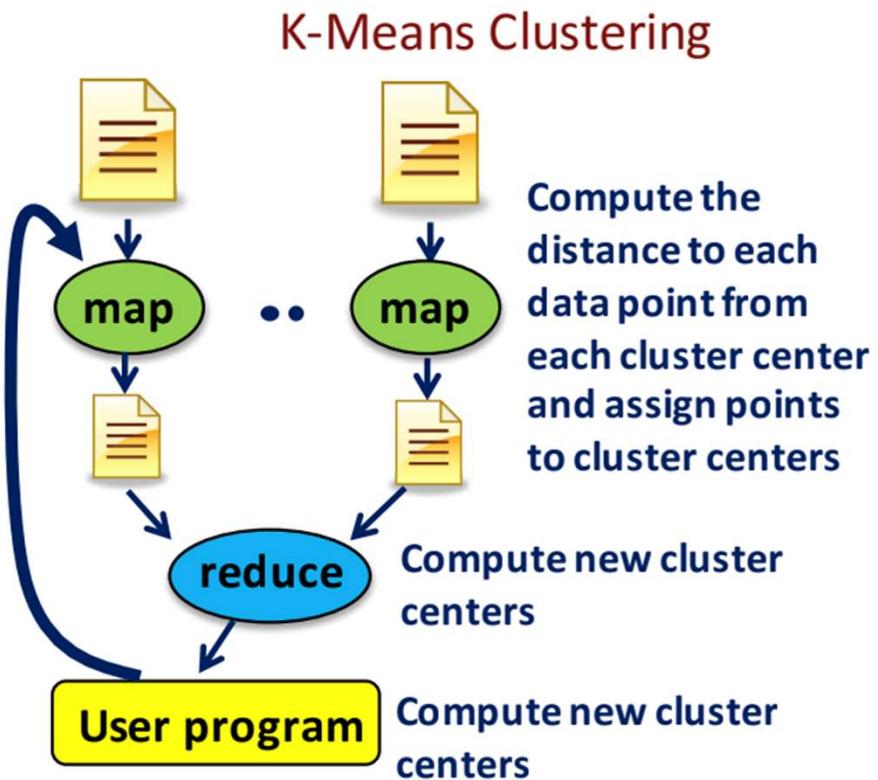
<http://shabal.in/visuals/kmeans/2.html>

# Iterative map reduce

- MapReduce is a one-pass computation
- Many applications, esp in ML and Data Mining areas, need to iteratively process data
- So, they need iterative execution of map reduce jobs
- An approach is to create a main program that calls the core map reduce with variable data
- Core program also checks for convergence
  - error bound (e.g. k-means clustering)
  - fixed iterations



# K-means as iterative map reduce



- The MapReduce program driver is responsible for repeating the steps via an iterative construct.
- Within each iteration map and reduce steps are called.
- Each map step reuses the result produced in previous reduce step.
  - e.g. k centers computed

<https://github.com/thomasjungblut/mapreduce-kmeans/tree/master/src/de/jungblut/clustering/mapreduce>

# K-Means Clustering by Iterative MapReduce

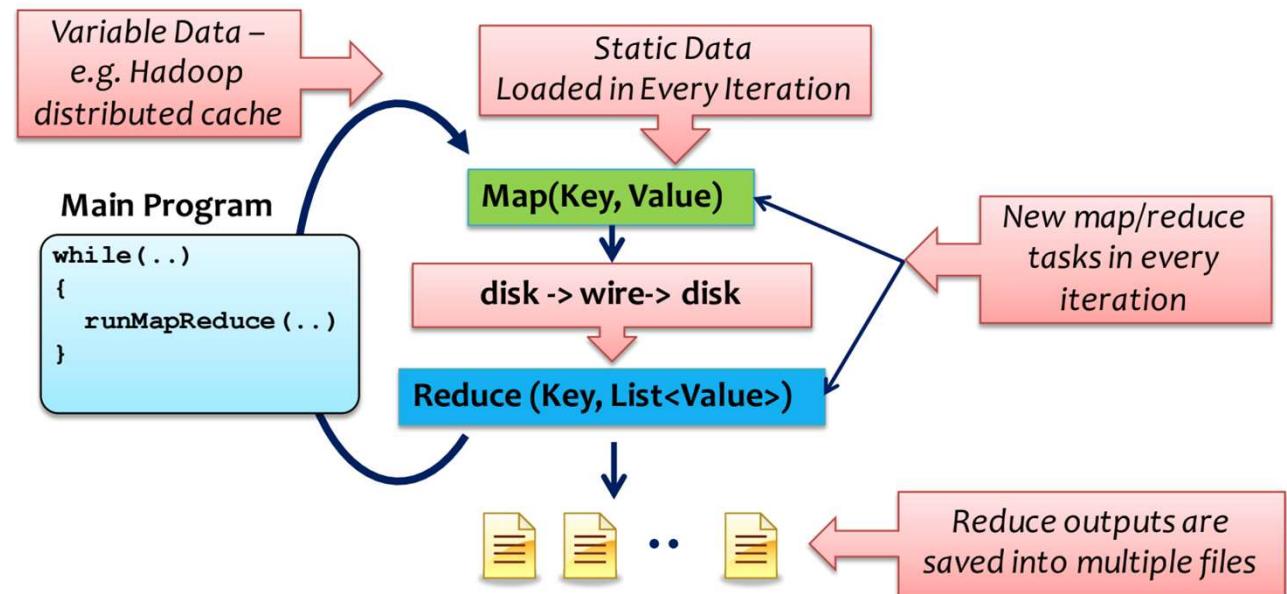
## Hands on demo

- 20Newsgroups folder - a set of around 20,000 postings to 20 different newsgroups with 1000 postings
- Convert all the newsgroup postings and turn them into bag-of-words vectors. – /data folder in HDFS
- Run a program that will choose an initial set of cluster centroids from the data – /clusters folder in HDFS (randomly sampled from the vectors)
- Run KMeans on Hadoop - hadoop jar MapRedKMeans.jar KMeans /data /clusters 3  
This will run 3 iterations of the KMeans algorithm on top of all documents in the 20\_newsgroups data set.  
This means that three separate MapReduce jobs will be run in sequence.
- The centroids produced at the end of:
  1. Iteration 1 will be put into the HDFS directory "/clusters1",
  2. Iteration 2 will be put into the HDFS directory "/clusters2",
  3. Iteration 3 will be put into the HDFS directory "/clusters3",

Reference: <https://cmj4.web.rice.edu/MapRedKMeans.html>

# Iterations using existing runtimes

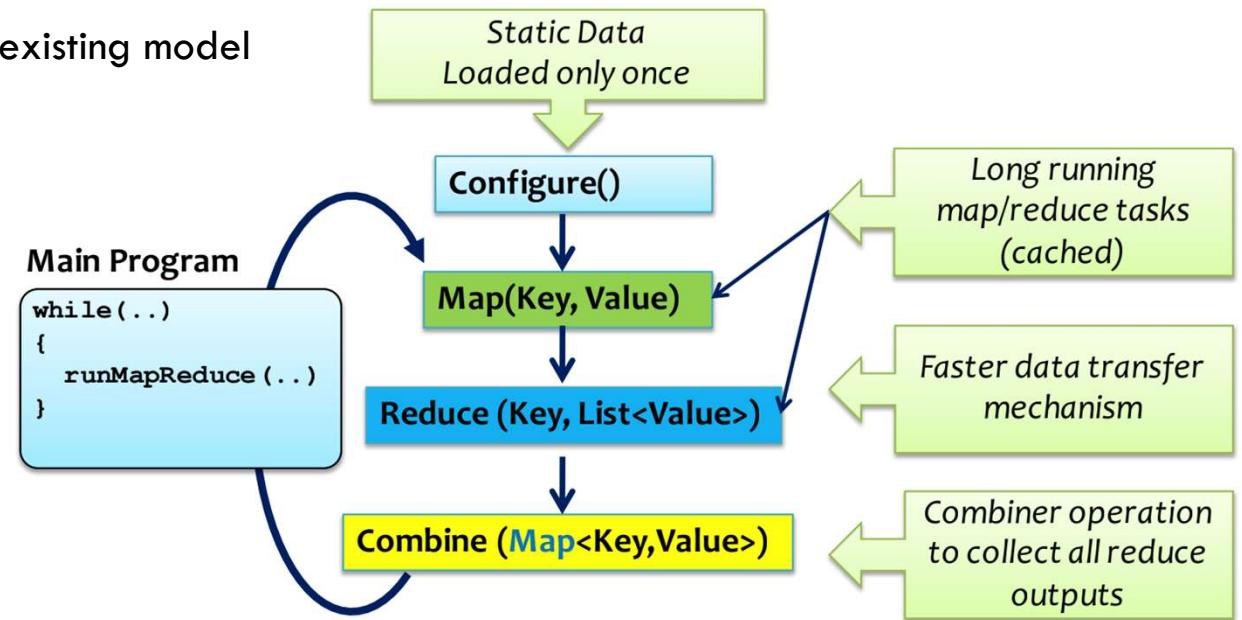
- Loop implemented on top of existing file-based single step map-reduce core
- Large overheads from
  - re-initialization of tasks
  - reloading of static data
  - communication and data transfers



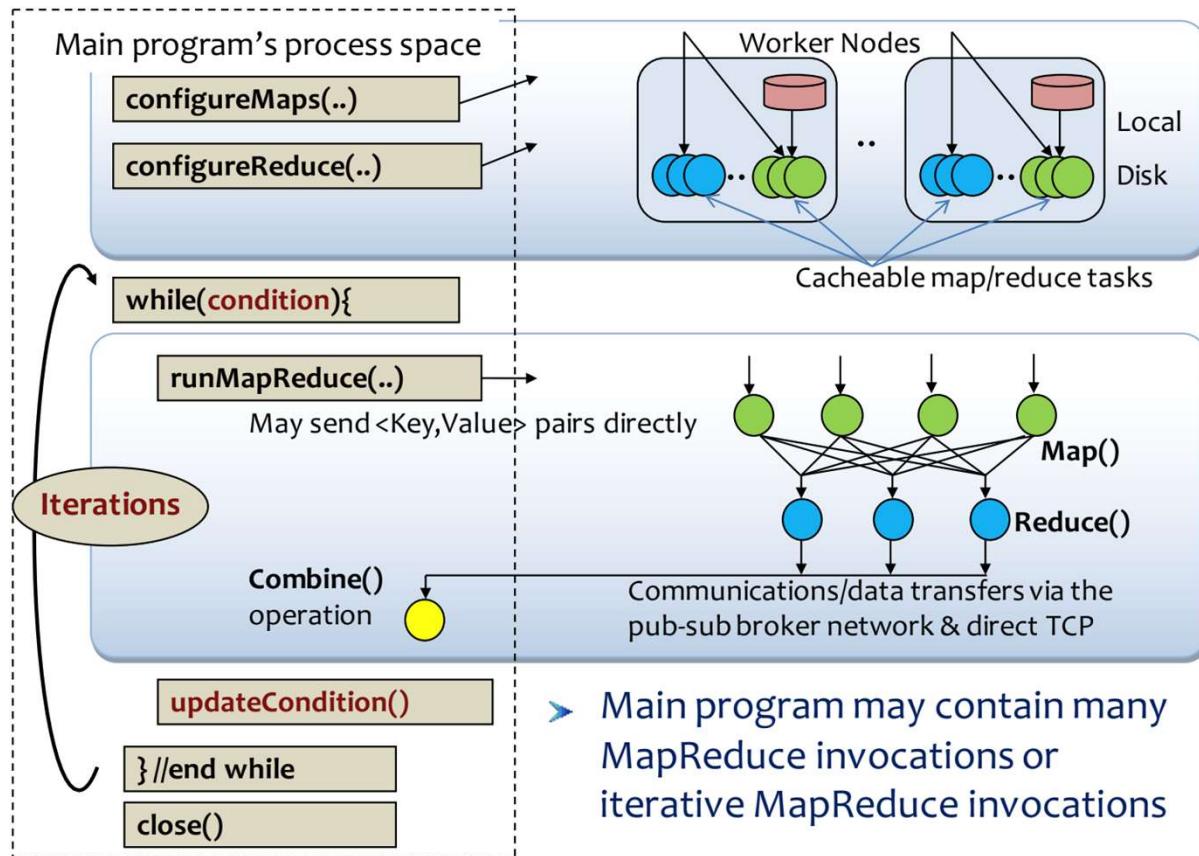
DistributedCache: <https://hadoop.apache.org/docs/r2.6.3/api/org/apache/hadoop/filecache/DistributedCache.html>

# MapReduce++ : Iterative MapReduce

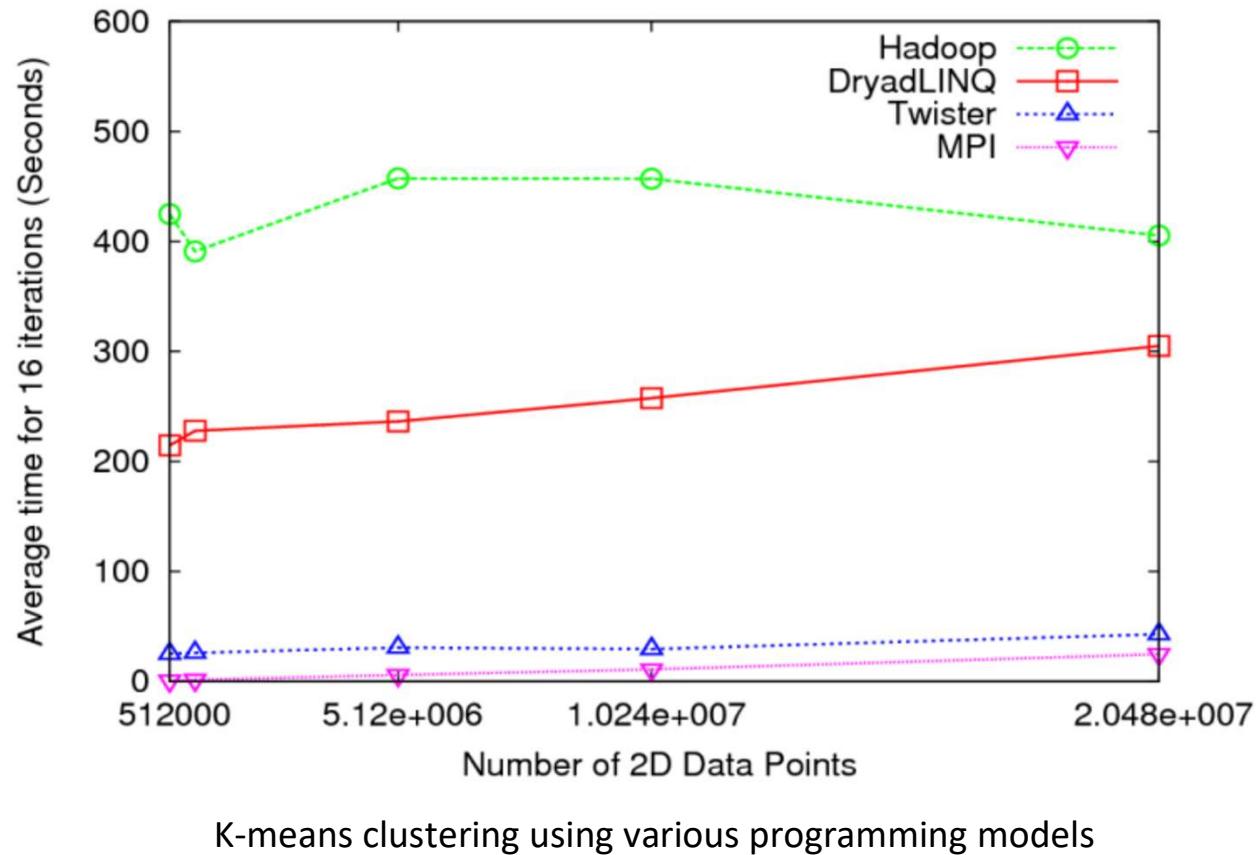
- Some optimizations are done on top of existing model
  - Static data loaded once
  - Cached tasks across invocations
  - Combine operations



# Example in Twister: Enables more APIs



# The optimisations indeed help

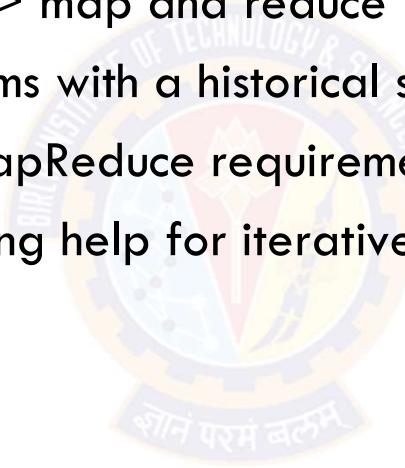


## Iterative MapReduce: Other options

- HaLoop
  - Modifies Hadoop scheduling to make it loop aware
  - Implements caches to avoid going to disk between iterations
  - Optional reading: Paper in [Proceedings of the VLDB Endowment](#) 3(1):285-296, Sep 2010
- Spark
  - Uses in-memory computing to speed up iterations
  - An in-memory structure called RDD : Resilient Distributed Dataset replaces files on disk
  - Ideal for iterative computations that reuse lot of data in each iteration

## Summary

- Different types of parallelism
- Data and tree parallelism —> map and reduce
- Basics of MapReduce programs with a historical sales data processing example
- Optimizations for iterative MapReduce requirements
- How does in-memory computing help for iterative MapReduce programming





**Next Session:**  
Hadoop MapReduce and YARN

