

Birla Institute of Technology & Science, Pilani
Work Integrated Learning Programmes Division
First Semester 2024-2025

Mid-Semester Test
(EC-2 Regular)

Course No. : DSECLZG524
 Course Title : Deep Learning
 Nature of Exam : Closed Book
 Weightage : 30%
 Duration : 2 Hours
 Date of Exam : 29-06-2025 (FN)

No. of Pages = 03
No. of Questions = 06

Note to Students:

1. Please follow all the *Instructions to Candidates* given on the cover page of the answer book.
 2. All parts of a question should be answered consecutively. Each answer should start from a fresh page.
 3. Assumptions made if any, should be stated clearly at the beginning of your answer.
-

Q.1 Consider the implementation of a perceptron for a NAND gate using bipolar representation (i.e., input and output values are either +1 or -1). Assume the following: **[5 Marks]**

- Initial weights: $w_0=w_1=w_2=0$, where w_0 is the bias and its input $x_0=1$
- Learning rate: $\eta=1$
- Activation function: $h=\text{sign}(z)$, where $z=w_0+w_1x_1+w_2x_2$
- Output: $y=h=\text{sign}(z)$
- Weight update rule: $\Delta w_i = \eta(t-y)x_i$, and $w_i \leftarrow w_i + \Delta w_i$

X ₁	X ₂	Target
-1	-1	1
-1	1	1
1	-1	1
1	1	-1

- A. Implement the perceptron learning algorithm for the above NAND gate using one full epoch (i.e., update the weights based on all four input combinations once).
- B. Show the computations for z , y , Δw , and updated weights for each input.
- C. Write the final weight vector (w_0 , w_1 , w_2) after the epoch.
- D. Verify that the perceptron correctly classifies all inputs with the final weights.

Note: Answer should be presented in tabular form similar to the example discussed in class.

Q2. Consider a binary classifier using a perceptron to classify a set of 2D points into two classes: red and blue. **[5 Marks]**

After plotting the data, you observe the following:

- Red and blue points cannot be separated by any straight line in 2D space.
- Multiple curved or combined boundaries could separate the classes successfully.

Answer the following:

- A. Why would a single-layer perceptron fail to classify this dataset correctly? [1 mark]
- B. How can a Multi-Layer Perceptron (MLP) overcome this limitation? Describe the role of hidden layers and non-linear activation functions. [2 Mark]
- C. For an XOR function over 6 binary input variables, how many perceptrons are required if you use a single hidden layer network? how many perceptrons are needed in a deep network? [2 Mark]

Note: Answer/Justify/Illustrate in no more than 30-50 words and a precise response w.r.t the given question only. Vague answers will be penalized.

Q3. A medical diagnostics company is developing two AI models: **[4 Marks]**

- Model A uses a neural network to detect if a patient has COVID-19 or not.
- Model B classifies patients into one of five disease categories based on symptoms.

Based on this context:

- A. Justify the use of a shallow neural network or a deep network for these models, respectively.
- B. Which loss function is appropriate for each system? Briefly justify your choices.
- C. What is the role of the SoftMax activation in Model B's output layer?
- D. Can SoftMax be used in System A? Explain why or why not.

Note: Answer/Justify/Illustrate in no more than 30-50 words and a precise response w.r.t the given question only. Vague answers will be penalized.

Q4. An industrial AI model is trained to predict energy consumption from sensor readings using a linear transformation. The model computes the following function: **[6 Marks]**

$$f(x, W) = \|W^\top x\|^2 = \sum_{j=1}^m (W^\top x)_j^2$$

Where:

- $x \in \mathbb{R}^2$ is the sensor input vector
- $W \in \mathbb{R}^{2 \times 2}$ is the weight matrix

Given:

$$W = \begin{bmatrix} 0.1 & 0.5 \\ -0.3 & 0.8 \end{bmatrix}, \quad x = \begin{bmatrix} 0.2 \\ 0.4 \end{bmatrix}$$

Answer the following:

- A. Draw a computational graph that clearly illustrates the operations from x, W to the scalar output $f(x, W)$.
- B. Perform the forward pass to compute the intermediate vector $q = W \cdot x$ and the final output f .
- C. Compute the gradient of the loss function for: $\nabla_q f, \nabla_W f, \nabla_x f$ (Show full matrix steps.)
- D. Explain in 2–3 lines how the computed gradients would be used to update W and x in a neural network training step.

Q5. Consider the following Python code snippet used to train a deep neural network on a multiclass classification task with 5 output classes, using the Keras library: **[5 Marks]**

```
model = Sequential()
model.add(Dense(64, input_dim=30, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='SGD', metrics=[ 'accuracy'])
model.fit(X_train, y_train, epochs=20, batch_size=16)

test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Loss: {test_loss}, Test Accuracy: {test_acc}")
```

Answer the following questions concisely:

- Describe the neural network architecture and the purpose of the activation functions used in each layer.
- Explain why categorical cross-entropy is appropriate for this setup. How would it change if labels were not one-hot encoded?
- What does a high training accuracy and low test accuracy suggest? Define overfitting and generalization in this context.
- Justify the use of the RMSprop optimizer. Mention one benefit and an alternative optimizer suitable for this task.
- Suggest two techniques to improve model generalization or training stability.

Note: (Answers must be concise within 30–50 words. Justifications are required for full credit.)

Q6. Answer the following questions: **[5 Marks]**

- You want to solve a classification task. You first train your network on 20 samples. Training converges, but the training loss is very high. You then decide to train this network on 10,000 examples. Is your approach to fixing the problem correct? If yes, explain the most likely results of training with 10,000 examples. If not, give a solution to this problem.
- A deep neural network for character recognition uses tanh in hidden layers and trains slowly with vanishing gradients. A switch to ReLU improves convergence.
 - Briefly explain why tanh caused issues.
 - Why does ReLU work better in this scenario?
- A fully connected neural network has 8 input neurons, two hidden layers with 10 and 6 neurons, and an output layer with 4 neurons. Biases are not included. Compute the total number of trainable weights.

Note: (Answers must be concise within 30–50 words. Justifications are required for full credit.)